

STOR 767 Spring 2019 Hw5: Computational Part

Due on 03/18/2019 in Class

Zhenghan Fang

Instruction.

- Please use **RMarkdown** to create a formatted report for the **Computational Part** of this homework.

Exercise 1

Define soft thresholding function and backfitting function.

```
soft.thresholding = function(z, lambda){
  if(abs(z) - lambda > 0){
    z.th = z / abs(z) * (abs(z) - lambda)
  }
  else{
    z.th = 0
  }
  return(z.th)
}

backfitting.LASSO = function(x, y, lambda, thresh=1e-20){
  N.pred = length(x[,1])
  n = length(x[,1])
  alpha = mean(y)
  beta = rep(0, N.pred)
  while(T){
    change = rep(0, N.pred)
    for(j in 1:N.pred){
      y.t = y - alpha - x[,-j] %*% beta[-j]
      beta.ls = (x[,j] %*% y.t) / n
      beta.lasso = soft.thresholding(beta.ls, lambda)
      change[j] = beta[j] - beta.lasso
      beta[j] = beta.lasso
    }
    if(sum(change^2) < thresh){
      break
    }
  }
  return(list("alpha" = alpha, "beta" = beta))
}
```

Perform backfitting. Left column: result of glmnet. Right column: result of my lasso.

```
x = matrix(rnorm(10000), 100, 100)
x.norm <- as.matrix(apply(x, 2, function(x) (x - mean(x))/sqrt(mean(x^2))))
y = rnorm(100)
lambda = 1/10
lasso.my = backfitting.LASSO(x, y, lambda)
lasso.glmnet = glmnet(x, y, alpha=1, lambda = 1/10, thresh=1e-20)
```

```
cbind(glmnet = cbind(glmnet = coef(lasso(glmnet), my = rbind(matrix(lasso.my$alpha), matrix(lasso.my$beta
```

```
## 101 x 2 sparse Matrix of class "dgCMatrix"
```

```
##              s0
```

```
## (Intercept) -0.110992175 -0.073078150
```

```
## V1          -0.027362148 -0.013308731
```

```
## V2          .            .
```

```
## V3          .            .
```

```
## V4          .            .
```

```
## V5          .            .
```

```
## V6          .            .
```

```
## V7          .            .
```

```
## V8          0.064776768  0.080489424
```

```
## V9          .            .
```

```
## V10         .            0.003502036
```

```
## V11         .            .
```

```
## V12         .            .
```

```
## V13         .            .
```

```
## V14         .            .
```

```
## V15         .            .
```

```
## V16         .            .
```

```
## V17         -0.011923815 -0.018826084
```

```
## V18         .            .
```

```
## V19         .            .
```

```
## V20         .            .
```

```
## V21         .            .
```

```
## V22         .            .
```

```
## V23         .            .
```

```
## V24         .            .
```

```
## V25         .            .
```

```
## V26         .            .
```

```
## V27         0.126737391  0.120936510
```

```
## V28         .            .
```

```
## V29         .            .
```

```
## V30         .            .
```

```
## V31         .            .
```

```
## V32         .            .
```

```
## V33         .            .
```

```
## V34         .            .
```

```
## V35         .            .
```

```
## V36         .            .
```

```
## V37         .            .
```

```
## V38         0.028759797  0.049783112
```

```
## V39         .            .
```

```
## V40         -0.017245754 -0.024655788
```

```
## V41         .            .
```

```
## V42         .            .
```

```
## V43         .            .
```

```
## V44         .            .
```

```
## V45         .            .
```

```
## V46         .            .
```

```
## V47         .            .
```

```
## V48         .            .
```

## V49	0.009007847	.
## V50	.	.
## V51	.	.
## V52	.	.
## V53	-0.180926676	-0.175595109
## V54	.	.
## V55	.	.
## V56	0.030402036	0.038209447
## V57	.	.
## V58	-0.017400607	-0.014057820
## V59	.	.
## V60	.	.
## V61	.	.
## V62	.	.
## V63	.	.
## V64	.	.
## V65	.	.
## V66	.	.
## V67	.	.
## V68	0.053150935	0.059852968
## V69	-0.045121408	-0.042835456
## V70	.	.
## V71	.	.
## V72	.	.
## V73	.	.
## V74	.	.
## V75	0.087606274	0.106623489
## V76	.	.
## V77	.	.
## V78	.	.
## V79	.	.
## V80	.	.
## V81	.	.
## V82	.	.
## V83	.	.
## V84	.	.
## V85	.	.
## V86	.	.
## V87	.	.
## V88	0.027219825	0.044945446
## V89	.	.
## V90	.	.
## V91	.	.
## V92	.	.
## V93	0.044784635	0.041944630
## V94	.	.
## V95	.	.
## V96	.	.
## V97	0.108262897	0.143128693
## V98	.	.
## V99	.	.
## V100	.	.

Exercise 2

Read data.

```
data.prost <- read.table('prostate.data.txt')
pred.name <- c("lcavol", "lweight", "age", "lbph", "svi", "lcp", "gleason", "pgg45")
resp.name <- "lpsa"
train.idx <- which(data.prost$train)
test.idx <- which(!data.prost$train)
pred.N <- length(pred.name)
formula.full <- as.formula(paste(resp.name, paste(pred.name, collapse=" + "), sep=" ~ "))
```

Standardize the predictors to unit variance and zero mean.

```
for (i in pred.name){
  t <- data.prost[, i]
  data.prost[, i] <- (t - mean(t)) / sqrt(var(t))
}
```

Split the data into training and test set.

```
data.train <- data.prost[train.idx, ]
data.test <- data.prost[test.idx, ]
```

Cross-validation

Define function “bs.cv” for cross validation. Inputs: data.train: training data; K.cv: number of folds.

```
bs.cv = function(data.train, K.cv){
  N.train = length(data.train[,1])
  cv.idx = sample(1:K.cv, N.train, replace=T)
  err.cv = c()
  for (k in 1:K.cv){
    bs = regsubsets(formula.full, data = data.train[cv.idx != k, ], nvmax=8)
    test.mat = model.matrix(formula.full, data = data.train[cv.idx == k, ])
    err.cv.temp = c()
    for (size.bs in 1:8){
      coefi = coef(bs, id = size.bs)
      pred = test.mat[, names(coefi)] %*% coefi
      err = (data.train[cv.idx == k, ]$lpsa - pred)^2
      err.cv.temp[size.bs] = mean(err)
    }
    err.cv = cbind(err.cv, err.cv.temp)
  }
  err.cv = rowMeans(err.cv)
  return(err.cv)
}
```

Perform 5-fold CV.

```
K.cv = 5
err.cv = bs.cv(data.train, K.cv)
```

The best subset size is

```
best.size.cv = which.min(err.cv)
best.size.cv
```

```
## [1] 7
```

The test error (MSE) is

```
bs = regsubsets(formula.full, data = data.train, nvmax=8)
test.mat = model.matrix(formula.full, data = data.test)
coefi= coef(bs, id = best.size.cv)
pred=test.mat[,names(coefi)]%*%coefi
err=(data.test$lpsa-pred)^2
mean(err)
```

```
## [1] 0.5165135
```

Perform 10-fold CV.

```
K.cv = 10
err.cv = bs.cv(data.train, K.cv)
```

The best subset size is

```
best.size.cv = which.min(err.cv)
best.size.cv
```

```
## [1] 7
```

The test error is

```
coefi= coef(bs, id = best.size.cv)
pred=test.mat[,names(coefi)]%*%coefi
err=(data.test$lpsa-pred)^2
mean(err)
```

```
## [1] 0.5165135
```

AIC and BIC

```
bs = regsubsets(formula.full, data = data.train, nvmax=8)
AIC.bs = c()
BIC.bs = c()
for(size.bs in 1:8){
  pred.name <- names(coef(bs,id=size.bs))
  formula.t <- as.formula(paste(resp.name, paste(pred.name[-1], collapse=" + "), sep=" ~ "))
  model.t <- lm(formula.t, data.train)
  AIC.bs[size.bs] = AIC(model.t)
  BIC.bs[size.bs] = BIC(model.t)
}
```

The best subset size from AIC is

```
best.size.AIC = which.min(AIC.bs)
best.size.AIC
```

```
## [1] 7
```

Test error is

```
coefi= coef(bs, id = best.size.AIC)
pred=test.mat[,names(coefi)]%*%coefi
err=(data.test$lpsa-pred)^2
mean(err)
```

```
## [1] 0.5165135
```

The best subset size from BIC is

```
best.size.BIC = which.min(BIC.bs)
best.size.BIC
```

```
## [1] 2
```

Test error is

```
cofi= coef(bs, id = best.size.BIC)
pred=test.mat[,names(cofi)]%*%cofi
err=(data.test$lpsa-pred)^2
mean(err)
```

```
## [1] 0.4924823
```

Bootstrap .632 estimator

```
bs = regsubsets(formula.full, data = data.train, nvmax=8)

err.632.bs = c()
theta.fit <- function(x,y){lsfit(x,y)}
theta.predict <- function(fit,x){
  cbind(1,x)%*%fit$coef
}
sq.err <- function(y,yhat) { (y-yhat)^2}
for(size.bs in 1:8){
  pred.name <- names(coef(bs,id=size.bs))
  x = data.train[,pred.name[-1]]
  y = data.train[,resp.name]
  results <- bootpred(x,y,200,theta.fit,theta.predict, err.meas=sq.err)
  err.632.bs[size.bs] = results[[3]]
}
```

The best subset size from bootstrap .632 estimator is

```
best.size.632 = which.min(err.632.bs)
best.size.632
```

```
## [1] 7
```

Test error is

```
cofi= coef(bs, id = best.size.632)
pred=test.mat[,names(cofi)]%*%cofi
err=(data.test$lpsa-pred)^2
mean(err)
```

```
## [1] 0.5165135
```

Exercise 3.

Load dataset.

```
SAheart = read.table('SAheart.data.txt', sep=",", header=T, row.names=1)
SAheart$chd = factor(SAheart$chd)
```

Use half of the dataset as training data.

```
train = sample(1:462, 231)
```

SVM with various kernels

SVM with linear kernel.

```
model.svm = svm(chd~., data=SAheart[train,], kernel = "linear")
summary(model.svm)
```

```
##
## Call:
## svm(formula = chd ~ ., data = SAheart[train, ], kernel = "linear")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##       cost:  1
##      gamma: 0.1
##
## Number of Support Vectors: 133
##
## ( 67 66 )
##
##
## Number of Classes: 2
##
## Levels:
## 0 1
```

Define error as number of misclassified samples/total number of samples.

Test error is

```
pred.svm = predict(model.svm, newdata = SAheart[-train, -10])
err = sum(pred.svm != SAheart$chd[-train]) / length(pred.svm)
err
```

```
## [1] 0.3073593
```

SVM with radial basis function kernel. Tune gamma by 10-fold cross validation.

```
model.svm = best.tune(svm, chd~., data=SAheart[train,], kernel = "radial", gamma = 0.05*2^(-1:3))
summary(model.svm)
```

```
##
## Call:
## best.tune(svm, chd ~ ., data = SAheart[train, ], kernel = "radial",
##      gamma = 0.05 * 2^(-1:3))
##
##
## Parameters:
```

```
## SVM-Type: C-classification
## SVM-Kernel: radial
## cost: 1
## gamma: 0.025 0.05 0.1 0.2 0.4
##
## Number of Support Vectors: 149
##
## ( 76 73 )
##
##
## Number of Classes: 2
##
## Levels:
## 0 1
```

Test error is

```
pred.svm = predict(model.svm, newdata = SAheart[-train, -10])
err = sum(pred.svm != SAheart$chd[-train] ) / length(pred.svm)
err
```

```
## [1] 0.2770563
```

SVM with sigmoid kernel. Tune gamma by 10-fold cross validation.

```
model.svm = best.tune(svm, chd~., data=SAheart[train,], kernel = "sigmoid", gamma = 0.05*2^(-1:3))
summary(model.svm)
```

```
##
## Call:
## best.tune(svm, chd ~ ., data = SAheart[train, ], kernel = "sigmoid",
## gamma = 0.05 * 2^(-1:3))
##
##
## Parameters:
## SVM-Type: C-classification
## SVM-Kernel: sigmoid
## cost: 1
## gamma: 0.025 0.05 0.1 0.2 0.4
## coef.0: 0
##
## Number of Support Vectors: 151
##
## ( 76 75 )
##
##
## Number of Classes: 2
##
## Levels:
## 0 1
```

Test error is

```
pred.svm = predict(model.svm, newdata = SAheart[-train, -10])
err = sum(pred.svm != SAheart$chd[-train] ) / length(pred.svm)
err
```

```
## [1] 0.2900433
```


LDA, QDA, and logistic regression

Perform LDA.

```
model = lda(formula=chd~., data=SAheart, subset=train)
model

## Call:
## lda(chd ~ ., data = SAheart, subset = train)
##
## Prior probabilities of groups:
##      0      1
## 0.6753247 0.3246753
##
## Group means:
##      sbp  tobacco      ldl adiposity famhistPresent      typea  obesity
## 0 137.2051 2.555192 4.286154 23.06404      0.2756410 52.91667 25.52622
## 1 145.8533 4.627333 5.682267 28.03667      0.5866667 54.52000 27.02280
##      alcohol      age
## 0 15.46923 38.67308
## 1 19.78573 50.78667
##
## Coefficients of linear discriminants:
##                      LD1
## sbp                0.001626368
## tobacco             0.040463910
## ldl                 0.225681772
## adiposity          -0.015040937
## famhistPresent     1.004167679
## typea              0.022398522
## obesity            0.009389706
## alcohol            0.002239492
## age                0.044142741
```

The test error of LDA is

```
pred = predict(object = model, newdata = SAheart[-train, ])
err = sum(pred$class != SAheart$chd[-train]) / length(pred$class)
err
```

```
## [1] 0.2597403
```

Perform QDA.

```
model = qda(formula=chd~., data=SAheart, subset=train)
model

## Call:
## qda(chd ~ ., data = SAheart, subset = train)
##
## Prior probabilities of groups:
##      0      1
## 0.6753247 0.3246753
##
## Group means:
##      sbp  tobacco      ldl adiposity famhistPresent      typea  obesity
## 0 137.2051 2.555192 4.286154 23.06404      0.2756410 52.91667 25.52622
## 1 145.8533 4.627333 5.682267 28.03667      0.5866667 54.52000 27.02280
```

```
##      alcohol      age
## 0 15.46923 38.67308
## 1 19.78573 50.78667
```

The test error of QDA is

```
pred = predict(object = model, newdata = SAheart[-train, ])
err = sum(pred$class != SAheart$chd[-train] ) / length(pred$class)
err
```

```
## [1] 0.3073593
```

Perform Logistic regression.

```
model = lrm(formula=chd~., data=SAheart, subset=train)
model
```

```
## Logistic Regression Model
```

```
##
## lrm(formula = chd ~ ., data = SAheart, subset = train)
##
##              Model Likelihood      Discrimination      Rank Discrim.
##              Ratio Test              Indexes              Indexes
## Obs          231  LR chi2      68.94  R2          0.360  C          0.813
## 0             156  d.f.          9    g           1.734  Dxy         0.626
## 1             75  Pr(> chi2) <0.0001 gr          5.660  gamma        0.626
## max |deriv| 1e-06                gp          0.279  tau-a         0.276
##                               Brier          0.161
##
##              Coef      S.E.    Wald Z Pr(>|Z|)
## Intercept    -7.7548 1.9187 -4.04 <0.0001
## sbp           0.0013 0.0078  0.17  0.8637
## tobacco      0.0404 0.0433  0.93  0.3516
## ldl          0.2439 0.0875  2.79  0.0053
## adiposity    -0.0120 0.0430 -0.28 0.7795
## famhist=Present 1.1164 0.3342  3.34 0.0008
## typea        0.0357 0.0179  1.99 0.0463
## obesity      0.0100 0.0602  0.17 0.8686
## alcohol      0.0051 0.0060  0.85 0.3956
## age          0.0662 0.0186  3.57 0.0004
##
```

The test error of Logistic regression is

```
pred = predict(object = model, newdata = SAheart[-train, ])
err = sum((pred>0)*1 != SAheart$chd[-train] ) / length(pred)
err
```

```
## [1] 0.2943723
```